

# Facing Off World Model Backbones: RNNs, Transformers, and S4

May 7, 2025

## Reference

Fei Deng, Junyeong Park, and Sungjin Ahn. Facing off world model backbones: Rnns, transformers, and s4, 2023. URL  
<https://arxiv.org/abs/2307.02064>.

# Linear State Space Models (SSMs)

- ▶ **Definition:** Maps 1-D input signal  $u(t)$  to 1-D output signal  $y(t)$
- ▶ **Formulation:**
  - ▶ **Continuous-time:**

$$\begin{aligned}\mathbf{s}'(t) &= \mathbf{A}\mathbf{s}(t) + \mathbf{B}u(t) \\ y(t) &= \mathbf{C}\mathbf{s}(t) + \mathbf{D}u(t)\end{aligned}$$

- ▶ **Discrete-time:**

$$\begin{aligned}\mathbf{s}_k &= \bar{\mathbf{A}}\mathbf{s}_{k-1} + \bar{\mathbf{B}}u_k \\ y_k &= \bar{\mathbf{C}}\mathbf{s}_k + \bar{\mathbf{D}}u_k\end{aligned}$$

- ▶ **Parallelizable SSMs (PSSMs):**
  - ▶ Combines benefits of RNNs and Transformers
  - ▶ Efficient autoregressive generation like RNNs
  - ▶ Parallelizable computation like Transformers
  - ▶ Computed via discrete convolution or parallel associative scan
- ▶ **PSSM Interface:**
  - ▶ **Single step:**  $\mathbf{y}_k, \mathbf{s}_k = \text{PSSM}(\mathbf{u}_k, \mathbf{s}_{k-1})$
  - ▶ **Parallel:**  $\mathbf{y}_{1:T}, \mathbf{s}_T = \text{PSSM}(\mathbf{u}_{1:T}, \mathbf{s}_0)$

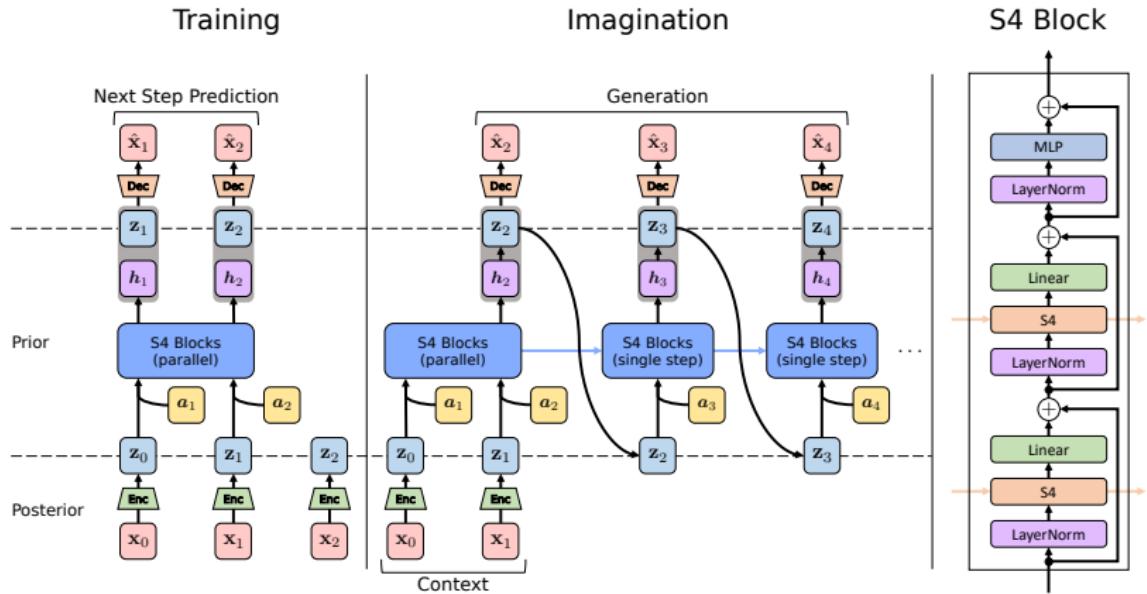
## The S4 Model

- ▶ **Core Idea:** SSMs with learnable parameters for sequence modeling
  - ▶ Matrices  $\mathbf{A}, \mathbf{B}, \mathbf{C}, \mathbf{D}$  and step size  $\Delta$  learned via gradient descent
  - ▶ Challenge: Expensive computation of  $\mathbf{A}$  powers
- ▶ **Key Innovation:** Diagonal Plus Low-Rank (DPLR) parameterization
  - ▶  $\mathbf{A} = \Lambda - \mathbf{P}\mathbf{P}^*$  with diagonal  $\Lambda$
  - ▶ HiPPO-LegS:  $\Lambda_{ii} = -(2i + 1)$ ,  $P_i = \sqrt{2i + 1}$
  - ▶ DPLR enables  $O(N)$  computation of powers and inverse via Woodbury identity
  - ▶ Negative diagonal ensures stability and uniform recall across timescales
- ▶ **Implementation:** HiPPO initialization, multiple SSMs for vector sequences
- ▶ **Extensions:** S4D (diagonal  $\mathbf{A}$ ), S5 (multi-input/output SSM)

## S4WM [Deng et al., 2023]: World Modeling with Structured State Space Models

- ▶ **Key Innovation:** First world-model framework compatible with any PSSM
  - ▶ Works with S4, S5, and their variants
  - ▶ Operates in compact latent space rather than raw pixels
  - ▶ Efficiently learns long-range dynamics in image sequences
- ▶ **Architecture Components:**
  - ▶ **Encoder:** Compresses raw observations into latent representations
  - ▶ **PSSM Backbone:** Models temporal dynamics in latent space
  - ▶ **Decoder:** Reconstructs observations from predicted latent states
- ▶ **Advantages:**
  - ▶ **Efficiency:** Lower computational cost than pixel-space modeling
  - ▶ **Long-range dependencies:** Captures complex temporal patterns
  - ▶ **Flexibility:** Compatible with various PSSM architectures
  - ▶ **Scalability:** Handles longer sequences than traditional approaches

# S4WM Architecture



**Figure:** S4WM, the first S4-based world model for improving long-term memory. S4WM efficiently models the long-range dependencies of environment dynamics in a compact latent space, using a stack of S4 blocks. This crucially allows fully parallelized training and fast recurrent latent imagination. S4WM is a general framework that is compatible with any parallelizable SSM including S5 and other S4 variants.

# Latent-Space Generative Process in S4WM

- ▶ **Stochastic Latent Variables:** Model future observations through latent space
  - ▶ Introduces stochastic latents  $\mathbf{z}_{0:T}$  to model observations  $\mathbf{x}_{1:T}$
  - ▶ Conditioned on initial observation  $\mathbf{x}_0$  and actions  $\mathbf{a}_{1:T}$
- ▶ **Generative Process:**

$$p(\mathbf{x}_{1:T} \mid \mathbf{x}_0, \mathbf{a}_{1:T}) = \int p(\mathbf{z}_0 \mid \mathbf{x}_0) \prod_{t=1}^T \underbrace{p(\mathbf{z}_t \mid \mathbf{z}_{<t}, \mathbf{a}_{\leq t})}_{\text{prior}} \underbrace{p(\mathbf{x}_t \mid \mathbf{z}_{\leq t}, \mathbf{a}_{\leq t})}_{\text{likelihood}} d\mathbf{z}_{0:T}.$$

- ▶ **Key Advantage:** Dimensionality Reduction
  - ▶ Factorization through latent variables avoids direct modeling of high-dimensional  $\mathbf{x}_t$
  - ▶ Enables efficient learning and inference in complex environments
  - ▶ PSSM backbone effectively captures temporal dependencies in compact latent space

## History Encoding via PSSM Blocks

- ▶ To capture dependencies in  $\{\mathbf{z}_{<t}, \mathbf{a}_{\leq t}\}$ , S4WM feeds

$$\mathbf{g}_t = \text{MLP}([\mathbf{z}_{t-1}; \mathbf{a}_t])$$

into a **stack** of PSSM blocks (e.g. S4 layers).

- ▶ These blocks compute in parallel during training and sequentially during rollout:

<b>(single step)</b>	$\mathbf{h}_t, \mathbf{s}_t = \text{PSSM\_Blocks}(\mathbf{g}_t, \mathbf{s}_{t-1})$
<b>(parallel)</b>	$\mathbf{h}_{1:T}, \mathbf{s}_T = \text{PSSM\_Blocks}(\mathbf{g}_{1:T}, \mathbf{s}_0)$

- ▶ Here  $\mathbf{h}_t$  is a fixed-size embedding summarizing all past latents and actions

# Probabilistic Components

- ▶ **Prior**

$$p(\mathbf{z}_t \mid \mathbf{z}_{\leq t}, \mathbf{a}_{\leq t}) = \text{MLP}(\mathbf{h}_t),$$

which outputs e.g. mean and variance for a Gaussian latent prior.

- ▶ **Likelihood**

$$p(\mathbf{x}_t \mid \mathbf{z}_{\leq t}, \mathbf{a}_{\leq t}) = \mathcal{N}(\hat{\mathbf{x}}_t, \mathbf{I}), \quad \hat{\mathbf{x}}_t = \text{Decoder}([\mathbf{h}_t; \mathbf{z}_t]).$$

This lets the model reconstruct pixels from the latent history

## Variational Training

- ▶ We introduce an CNN **encoder** that maps each  $\mathbf{x}_t$  to a posterior over  $\mathbf{z}_t$ :

$$q(\mathbf{z}_{0:T} \mid \mathbf{x}_{0:T}, \mathbf{a}_{1:T}) = \prod_{t=0}^T q(\mathbf{z}_t \mid \mathbf{x}_t), \quad q(\mathbf{z}_0 \mid \mathbf{x}_0) = p(\mathbf{z}_0 \mid \mathbf{x}_0).$$

- ▶ The **ELBO** objective is

$$\begin{aligned} \log p(\mathbf{x}_{1:T} \mid \mathbf{x}_0, \mathbf{a}_{1:T}) &\geq \\ \mathbb{E}_q \left[ \sum_{t=1}^T \log p(\mathbf{x}_t \mid \mathbf{z}_{\leq t}, \mathbf{a}_{\leq t}) - \text{KL}(q(\mathbf{z}_t \mid \mathbf{x}_t) \parallel p(\mathbf{z}_t \mid \mathbf{z}_{<t}, \mathbf{a}_{\leq t})) \right]. \end{aligned}$$

- ▶ All  $\mathbf{z}_t$  samples and PSSM computations run in parallel during training

# Training Algorithm

---

**Algorithm 1:** S4WM Training

---

**Input:**  $(\mathbf{x}_0, \mathbf{a}_1, \mathbf{x}_1, \dots, \mathbf{a}_T, \mathbf{x}_T)$

▷ Obtain posterior latents

**for** time step  $t = 0, \dots, T$  **parallel do**  
  └  $\mathbf{z}_t \sim q(\mathbf{z}_t | \mathbf{x}_t) = \text{Encoder}(\mathbf{x}_t)$

▷ Prepare inputs to S4 blocks

**for** time step  $t = 1, \dots, T$  **parallel do**  
  └  $\mathbf{g}_t = \text{MLP}([\mathbf{z}_{t-1}; \mathbf{a}_t])$

▷ Encode history by S4 blocks

$\mathbf{h}_{1:T}, \mathbf{s}_T = \text{S4Blocks}(\mathbf{g}_{1:T}, \mathbf{s}_0)$

▷ Compute prior and decode posterior latents

**for** time step  $t = 1, \dots, T$  **parallel do**  
  └  $p(\mathbf{z}_t | \mathbf{z}_{<t}, \mathbf{a}_{\leq t}) = \text{MLP}(\mathbf{h}_t)$   
    └  $\hat{\mathbf{x}}_t = \text{Decoder}([\mathbf{h}_t; \mathbf{z}_t])$

---

Compute objective by ELBO Equation

---

---

**Algorithm 2:** S4WM Imagination

---

**Input:** context  $(\mathbf{x}_{0:C}, \mathbf{a}_{1:C})$ , query  $\mathbf{a}_{C+1:T}$

▷ Encode context

**for** *context step*  $t = 0, \dots, C$  **parallel do**

$\mathbf{z}_t \sim q(\mathbf{z}_t | \mathbf{x}_t) = \text{Encoder}(\mathbf{x}_t)$

$\mathbf{g}_{t+1} = \text{MLP}([\mathbf{z}_t; \mathbf{a}_{t+1}])$

$\mathbf{h}_{1:C+1}, \mathbf{s}_{C+1} = \text{S4Blocks}(\mathbf{g}_{1:C+1}, \mathbf{s}_0)$

▷ Imagine in latent space

$\mathbf{z}_{C+1} \sim p(\mathbf{z}_{C+1} | \mathbf{z}_{0:C}, \mathbf{a}_{1:C+1}) = \text{MLP}(\mathbf{h}_{C+1})$

**for** *query step*  $t = C + 2, \dots, T$  **do**

$\mathbf{g}_t = \text{MLP}([\mathbf{z}_{t-1}; \mathbf{a}_t])$

$\mathbf{h}_t, \mathbf{s}_t = \text{S4Blocks}(\mathbf{g}_t, \mathbf{s}_{t-1})$

$\mathbf{z}_t \sim p(\mathbf{z}_t | \mathbf{z}_{<t}, \mathbf{a}_{\leq t}) = \text{MLP}(\mathbf{h}_t)$

▷ Decode imagined latents

**for** *query step*  $t = C + 1, \dots, T$  **parallel do**

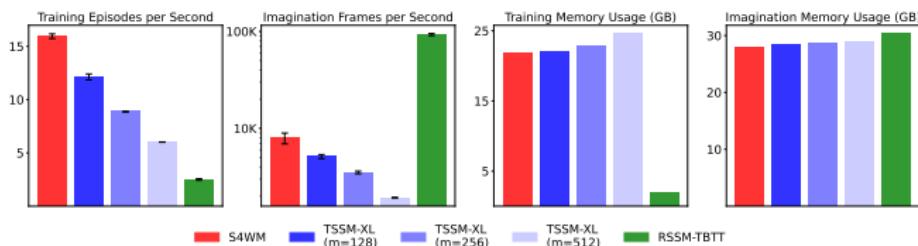
$\hat{\mathbf{x}}_t = \text{Decoder}([\mathbf{h}_t; \mathbf{z}_t])$

# Environments for Evaluating World Model Memory Capabilities

- ▶ **Evaluation Focus:** Memory capabilities in world models
  - ▶ Long-term imagination
  - ▶ Context-dependent recall
  - ▶ Reward prediction
  - ▶ Memory-based reasoning
- ▶ **Environment Design:**
  - ▶ 3D Memory Maze and 2D MiniGrid environments with partial observability
  - ▶ Each environment targets specific memory capability
  - ▶ Deterministic environments with moderate visual complexity
- ▶ **Evaluation Protocol:**
  - ▶ World model observes context phase
  - ▶ Evaluated on imagination given query phase actions
  - ▶ Mean squared error (MSE) as primary metric
  - ▶ Evaluation on unseen episodes from same policy distribution
- ▶ **Advantages over Prior Work:**
  - ▶ Provides deeper insights than final agent performance
  - ▶ Isolates world model capabilities from policy learning
  - ▶ Enables focused assessment of specific memory aspects
  - ▶ Paves way for improved model-based agents with better memory

# World Model Baselines and Performance Comparison

- ▶ **Baseline Models:**
  - ▶ **RSSM-TBTT:** RNN-based world model with truncated backpropagation
  - ▶ **TSSM-XL:** Transformer-based model using Transformer-XL for longer sequences
  - ▶ **S4WM:** Structured State Space Model-based world model
- ▶ **Computational Efficiency:**
  - ▶ **Training Speed:** S4WM and TSSM-XL train faster due to parallel computation
  - ▶ **Memory Usage:** RSSM-TBTT is most memory-efficient during training
  - ▶ **Imagination Throughput:** RSSM-TBTT achieves  $\sim 10\times$  faster inference

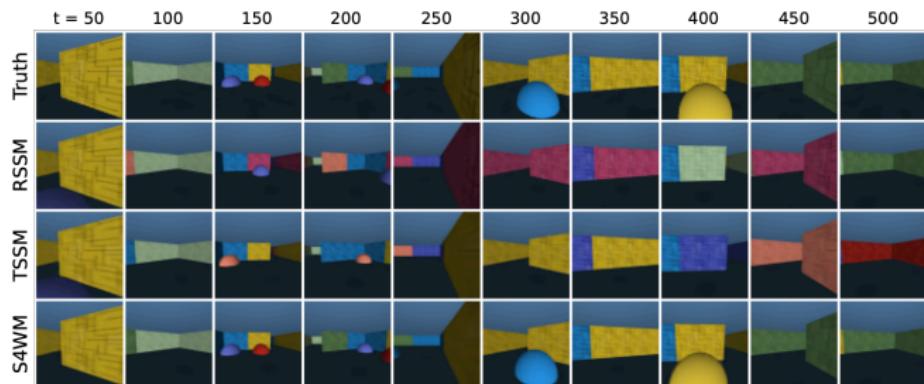


**Figure:** Comparison of speed and memory usage during training and imagination. S4WM is the fastest to train, while RSSM-TBTT is the most memory-efficient during training and has the highest throughput during imagination.

# Long-Term Imagination

	Two Rooms (301   200)		Four Rooms (501   500)		Ten Rooms (1101   900)	
	Recon. MSE ( $\downarrow$ )	Gen. MSE ( $\downarrow$ )	Recon. MSE ( $\downarrow$ )	Gen. MSE ( $\downarrow$ )	Recon. MSE ( $\downarrow$ )	Gen. MSE ( $\downarrow$ )
RSSM-TBTT	<b>1.7</b>	62.2	<b>1.5</b>	219.4	<b>1.5</b>	323.1
TSSM-XL	2.5	62.9	2.4	224.4	2.6	360.4
S4WM	1.8	<b>27.3</b>	1.7	<b>44.0</b>	1.8	<b>224.4</b>

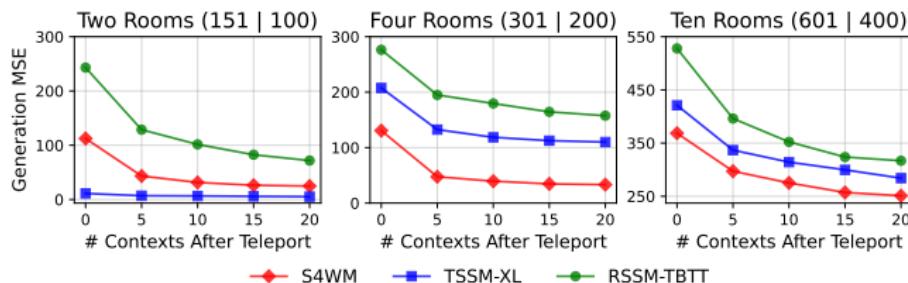
**Table:** Long-term imagination evaluation. Format: (context | query steps). All models show good reconstruction, but S4WM excels at generation up to 500 steps. Ten Rooms remains challenging for all models.



**Figure:** Long-term imagination in Four Rooms. RSSM-TBTT and TSSM-XL show significant errors in walls, objects, and positions, while S4WM generates more accurately with only minor position errors.

## Context-Dependent Recall

- ▶ **Teleport Experiment:** Agent observes initial context, then is teleported to random point in history and must recall subsequent events using the same action sequence
- ▶ **Evaluates:** Models' ability to access specific memories from arbitrary points in history



**Figure:** Context-dependent recall in teleport environments (format: context | query steps). TSSM-XL excels with short contexts, requiring no additional observations, while S4WM performs best with longer context phases.

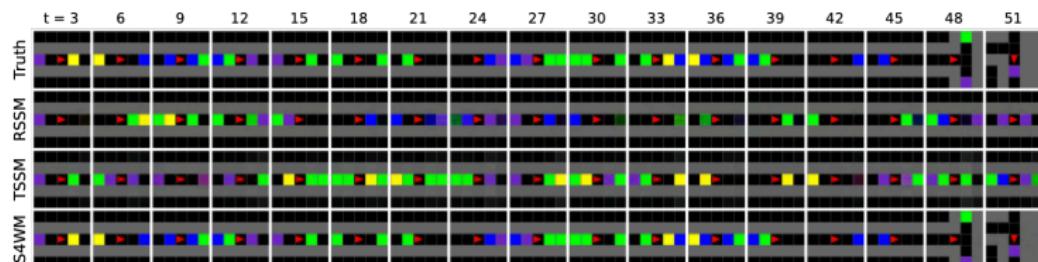
## Reward Prediction

- ▶ To facilitate policy learning within imagination, world models need to accurately predict the rewards
- ▶ We evaluate reward prediction accuracy over long time horizons using the visually simpler 2D MiniGrid environment
- ▶ **Distracting Memory environment:**
  - ▶ More challenging than original MiniGrid Memory environment
  - ▶ Distractors of random colors placed in hallway
  - ▶ Reward of 1 given if square reached matches color of square in left room
  - ▶ Model must ignore distractors while tracking agent's position
- ▶ **Results:**
  - ▶ Only S4WM accurately predicts rewards within imagination
  - ▶ TSSM-XL succeeds with full sequence but fails in imagination
  - ▶ RSSM-TBTT completely fails (close to random guessing)
  - ▶ Failures mainly due to inability to track agent's position

## Reward Prediction (cont.)

**Table:** Reward prediction accuracy in the Distracting Memory environments. Each environment is labeled with (context steps | query steps). S4WM succeeds in all environments. TSSM-XL has limited success when observing the full sequence, but fails to predict rewards within imagination. RSSM-TBTT completely fails.

	Width = 100 (199   51)		Width = 200 (399   101)		Width = 400 (799   201)	
	Inference Accuracy (↑)	Imagination Accuracy (↑)	Inference Accuracy (↑)	Imagination Accuracy (↑)	Inference Accuracy (↑)	Imagination Accuracy (↑)
RSSM-TBTT	47.9%	49.6%	48.7%	48.4%	50.4%	52.2%
TSSM-XL	<b>100.0%</b>	51.2%	99.9%	51.3%	50.4%	51.0%
S4WM	<b>100.0%</b>	<b>100.0%</b>	<b>100.0%</b>	<b>100.0%</b>	<b>100.0%</b>	<b>100.0%</b>



**Figure:** Imagination in the Distracting Memory environment of width 100. RSSM-TBTT and TSSM-XL fail to keep track of the agent's position, leading to inaccurate reward prediction within imagination.