

Simulus, SGF

July 2, 2025

Reference

Lior Cohen, Kaixin Wang, Bingyi Kang, Uri Gadot, and Shie Mannor.
Uncovering untapped potential in sample-efficient world model agents, 2025.
URL <https://arxiv.org/abs/2502.11537>.

Uncovering Untapped Potential in Sample-Efficient World Model Agents (Simulus)

Simulus [Cohen et al., 2025]

- ▶ **Simulus** builds on top of REM and following the same \mathcal{V} - \mathcal{M} - \mathcal{C} structure.
- ▶ Tokenizer \mathcal{V} encodes raw observations o_t and actions a_t into *fixed-length* sequence.
- ▶ World model \mathcal{M} parameterized by θ , would embed and stream into a *RetNet* world model f_θ . (POP lets the RetNet predict the whole next observation in one go instead of token-by-token.)
- ▶ An ensemble of 4 identical heads predicts the same tokens. Their Jensen-Shannon disagreement is the intrinsic reward i_t .
- ▶ A LSTM Actor-Critic \mathcal{C} is optimized with:

$$\bar{r}_t = \alpha_{\text{ext}} \hat{r}_t + \alpha_{\text{int}} i_t.$$

- ▶ Prioritized WM replay: world-model batches are 30% high-loss frames, 70% uniform.

Tokenizer \mathcal{V}

- ▶ Each modality would have its own encoder/decoder pair.

| Modality | Encoder | Tokens per obs | Vocab |
|----------------------|---|-------------------|-------------------|
| Image 64×64 | 3-level VQ-VAE | $8 \times 8 = 64$ | 512 |
| Vector (cont.) | Scalar \rightarrow 125-bin quantiser \pm symlog | len(vector) | 125 |
| Categorical | identity (already integer) | len(cat) | native |
| 2-D grid | flatten, embed & average | mn | sizes per channel |

Table: Tokenizer modalities and their encoding specifications

World Model \mathcal{M}

World Model input:

- ▶ Tokenizer converts raw observations o_t into tokens $z_t = (z_{t,1}, \dots, z_{t,K})$.
- ▶ Each token $z_{t,j}$ is a small integer ($0 \dots \text{vocab_size} - 1$)
- ▶ Tokens are mapped to fixed-size embeddings $e_{t,j} \in \mathbb{R}^d$
- ▶ Concatenate K vectors to form observation block $E_t = (e_{t,1}, \dots, e_{t,K})$
- ▶ Action tokens A_t are appended to E_t , forming input pair $\underbrace{E_t}_{\text{observation}}, \underbrace{A_t}_{\text{action tokens}}$.
- ▶ The world model's input stream follows the pattern: $E_1 A_1 E_2 A_2 E_3 A_3 \dots$

World Model architecture:

- ▶ **RetNet Architecture:** Transformer-like network where expensive self-attention is replaced by cheaper **Retention** operation
- ▶ **Sequential Operation:** At each step consumes one full (E, A) block and updates hidden state:

$$(h_t, x_t) = f_\theta(h_{t-1}, (E_t, A_t))$$

where h_t : recurrent state, x_t : prediction of next observation.

- ▶ **POP:** Uses learnable query sequence \mathbf{U} to predict all K tokens in parallel
- ▶ **POP Implementation:**
 1. After processing block t , hold hidden state h_t
 2. Call RetNet with learnable \mathbf{U} : $(_, y_t^u) = f_\theta(h_t, \mathbf{U})$
 3. Feed $y_t^u \in \mathbb{R}^{K \times d}$ rows into MLP heads to get $p_\theta(\hat{z}_{t+1,j} | y_{t,j}^u)$
- ▶ **Result:** Parallel K predictions enable $K \times$ faster generation

Intrinsic Reward Signal

Motivation: Extrinsic reward is sparse early in training. Model-based agents can measure their own ignorance - prediction disagreements indicate where more data is needed.

Implementation:

- ▶ **4 Independent Heads:** $p_{\phi_1}, p_{\phi_2}, p_{\phi_3}, p_{\phi_4}$ all predict from same RetNet
- ▶ **Jensen-Shannon Divergence (JSD)** measures disagreement:

$$u_j = H\left(\frac{1}{4} \sum_{i=1}^4 p_{\phi_i}\right) - \frac{1}{4} \sum_{i=1}^4 H(p_{\phi_i})$$

- ▶ **Average over tokens:** $i_t = \frac{1}{K} \sum_{j=1}^K u_j$
- ▶ **Combined reward:** $\bar{r}_t = \alpha_{\text{ext}} \hat{r}_t + \alpha_{\text{int}} i_t$

Architecture:

- ▶ All heads see **stop-gradient** inputs for JSD computation
- ▶ JSD bounded in $[0, \log V]$ - mixes well with any reward scale
- ▶ High JSD frames get high intrinsic reward and prioritized replay
- ▶ Creates feedback loop: explore → learn → reduce uncertainty → shift exploration

Prioritized Replay for World Model

Implementation:

- ▶ Store latest obs-loss with each transition
- ▶ WM batch sampling: 70% uniform, 30% softmax(loss)
- ▶ New frames: high dummy loss for guaranteed sampling

Benefits:

- ▶ Focus on hard-to-predict transitions
- ▶ New experiences always sampled
- ▶ Balanced learning via uniform component

Controller \mathcal{C}

Overview: Outputs stochastic policy $\pi_\psi(a_t \mid \tau_t)$ and value estimate $V_\psi(\tau_t)$

Recurrent Backbone:

- ▶ Single-layer LSTM
- ▶ $h_t, c_t = \text{LSTM}(v_t, h_{t-1}, c_{t-1})$
- ▶ Constitutes majority of controller parameters ψ

Input Pipeline:

- ▶ **Modality Encoders:** Each token type $z_t^{(i)}$ gets encoded via $E^{(i)} : \{0, \dots, V_i - 1\}^{K_i} \rightarrow \mathbb{R}^{d_i}$
- ▶ **Fusion MLP:** $v_t = g_{\text{fuse}}([v_t^{(0)} \| v_t^{(1)} \| \dots]) \in \mathbb{R}^{d_{\text{fuse}}}$

Training:

- ▶ Generate imagined trajectory $\hat{\tau} = (z_1, a_1, \bar{r}_1, \dots, z_H, a_H, \bar{r}_H)$
- ▶ λ -returns + REINFORCE with value baseline + entropy regularization like Dreamer.

Results




Figure: Simulus results

Simple, Good, Fast: Self-Supervised World Models Free of Baggage (SGF)

Aim: Strip world-models to bare essentials - can we do well without RNNs/Transformers, discrete latents, or pixel reconstructions?

Core Ingredients:

- ▶ Self-supervised representations (VICReg-style)
- ▶ Two losses: temporal consistency + information maximization
- ▶ Frame + action stacking (captures short-term dependencies)
- ▶ Strong image augmentations (injects stochasticity)
- ▶ Deterministic, feed-forward dynamics (MLP predicts Δ -latent, reward, terminal)

Why it works:

- ▶ Learned latents are informative and locally smooth → simple dynamics suffices
- ▶ Augmentations stand in for stochastic modelling

Ingredients Leading to Simplicity

Stacking instead of memory:

- ▶ Traditional: RNNs/attention for long-term dependencies
- ▶ Their approach: *frame and action stacking*
- ▶ Stack m recent frames → captures velocity, partial observability
- ▶ Stack m recent actions → handles action delays
- ▶ Much faster than recurrent/attention mechanisms

Augmentations instead of stochasticity:

- ▶ Prior models: stochastic even in deterministic POMDPs
- ▶ Our approach: *data augmentation* for stochasticity
- ▶ Random augmentations during training → robustness
- ▶ Avoids computational overhead of stochastic predictions
- ▶ Similar to successful model-free approaches (DrQ, RAD)

Key insight: Simple alternatives achieve similar performance with much less complexity

Model Structure




Figure: Model Structure

Representation Learning

- Given transition $(\mathbf{o}, \mathbf{a}, \mathbf{o}', r, e)$, apply random augmentations $t, t' \sim \mathcal{T}$ to get $\tilde{\mathbf{o}} = t(\mathbf{o})$ and $\tilde{\mathbf{o}}' = t'(\mathbf{o}')$
- Encoder f_{enc} computes representations $\tilde{\mathbf{y}} = f_{\text{enc}}(\tilde{\mathbf{o}})$ and $\tilde{\mathbf{y}}' = f_{\text{enc}}(\tilde{\mathbf{o}}')$
- Projector f_{proj} computes embeddings $\tilde{\mathbf{z}} = f_{\text{proj}}(\tilde{\mathbf{y}})$ and $\tilde{\mathbf{z}}' = f_{\text{proj}}(\tilde{\mathbf{y}}')$
- Action-conditioned predictor f_{pred} predicts $\tilde{\mathbf{z}}' = f_{\text{pred}}(\tilde{\mathbf{z}}, \mathbf{a})$

Loss Function:

$$\mathcal{L}_{\text{Repr.}}(\theta) = \mathbb{E}_{\tau} \left[\underbrace{\frac{\eta}{D} \|f_{\text{pred}}(\tilde{\mathbf{z}}, \mathbf{a}) - \tilde{\mathbf{z}}'\|_2^2}_{\text{Temporal Consistency}} + \underbrace{\text{VC}(\tilde{\mathbf{Z}}) + \text{VC}(\tilde{\mathbf{Z}}')}_{\text{Information Maximization}} \right]$$

Variance-Covariance Regularization:

$$\text{VC}(\mathbf{Z}) = \frac{1}{D} \sum_{j=1}^D \left[\underbrace{\rho \max \left(0, 1 - \sqrt{\text{Cov}(\mathbf{Z})_{j,j} + \varepsilon} \right)}_{\text{Variance}} + \underbrace{\nu \sum_{k \neq j} \text{Cov}(\mathbf{Z})_{j,k}^2}_{\text{Covariance}} \right]$$

Where \mathbf{Z} is the set of all embeddings in the batch.

Variance regularization keeps the standard deviation of each embedding feature across the batch above 1 using a hinge loss. Covariance regularization decorrelates the embedding features by attracting their covariances towards zero

Dynamics Model

Maximum likelihood estimation:

$$\mathcal{L}_{\text{Dyn.}}(\theta) = \mathbb{E}_\tau \left[-\underbrace{\log p_\theta(\text{sg}(\mathbf{y}') | \text{sg}(\mathbf{y}), \mathbf{a})}_{\text{Transition Distribution}} - \underbrace{\log p_\theta(r | \tilde{\mathbf{y}}, \mathbf{a}, \tilde{\mathbf{y}'})}_{\text{Reward Distribution}} - \underbrace{\log p_\theta(e | \tilde{\mathbf{y}}, \mathbf{a}, \tilde{\mathbf{y}'})}_{\text{Terminal Distribution}} \right]$$

Key design choices:

- ▶ Stop-gradient $\text{sg}(\cdot)$ on representations in transition loss
- ▶ Prevents moving targets from representation model updates
- ▶ Rewards/terminals provide stable POMDP signals
- ▶ Learn transitions with non-augmented observations: $\mathbf{y} = f_{\text{enc}}(\mathbf{o})$, $\mathbf{y}' = f_{\text{enc}}(\mathbf{o}')$

Results




Figure: Score and runtime comparison in the Atari 100k benchmark. SPR is model-free, EfficientZero performs lookahead.

Limitations

1. Limited to MDPs, not POMDPs:

- ▶ Only works with **deterministic MDPs**
- ▶ Cannot handle **non-deterministic POMDPs** because:
 - ▶ Transition distribution needs to be **stochastic**
 - ▶ Predictor must handle **uncertainty** between observations
 - ▶ Both networks need **stochastic predictions** (mean/variance or Gaussian mixtures)
- ▶ Limited to **short-term dependencies** (avoided RNN/Transformers)
- ▶ This might explain why we didn't reach **state-of-the-art performance**

2. VICReg Image Requirement:

- ▶ VICReg currently requires **image observations**
- ▶ Could be applied to **other modalities** if reasonable augmentations are available
- ▶ Could combine SGF with **other self-supervised learning methods**