

Curious Replay for Model-based Adaptation

July 2, 2025

Reference

Isaac Kauvar, Chris Doyle, Linqi Zhou, and Nick Haber. Curious replay for model-based adaptation, 2023. URL <https://arxiv.org/abs/2306.15934>.

Curious Replay for Model-based Adaptation ([Kauvar et al., 2023])

Curious Replay Algorithm

```
1: Input: Replay buffer  $R$  that uses a SumTree structure to store the priority  
    $p_i$  of each transition  
2: Hyperparameters:  $c, \beta, \alpha, \epsilon$ , environment steps per train step  $L$ , batch  
   size  $B$ , maximum priority  $p_{\text{MAX}}$   
3: for iteration 1, 2, ... do  
4:   Collect  $L$  transitions  $(x_t, a_t, r_t, x_{t+1})$  with policy  
5:   Add transitions to replay buffer  $R$ , each with priority  $p_i \leftarrow p_{\text{MAX}}$  and visit  
   count  $v_i \leftarrow 0$   
6:   Sample batch of  $B$  transitions from  $R$  using probability for selecting  
   transition  $i$  as  $p_i / \sum_{j=1}^{|R|} p_j$   
7:   Train world model and policy using batch, and cache loss  $\mathcal{L}_i$  for each  
   transition in batch  
8:   for transition  $i$  in batch do  
9:      $p_i \leftarrow c\beta^{v_i} + (|\mathcal{L}_i| + \epsilon)^\alpha$   
10:     $v_i \leftarrow v_i + 1$   
11:   end for  
12: end for
```

Curious Replay: Motivation

- ▶ **Challenge:** After an environment change, the world model is suddenly inaccurate on new data.
- ▶ **Goal:** Adapt quickly by focusing learning on the most relevant transitions.
- ▶ **Solution: Curious Replay** prioritizes replay buffer sampling to:
 - ▶ Fix transitions the model predicts poorly *right now*
 - ▶ Ensure new, rarely trained transitions are not neglected

Curious Replay: The Two Key Signals

1. **Model Error:** Use the world model's prediction loss \mathcal{L}_i for each transition i .
 - ▶ Higher $|\mathcal{L}_i|$ means the model is more wrong → sample it more.
2. **Under-Replay:** Track a visit count v_i for each transition.
 - ▶ Fewer replays (lower v_i) → boost its priority to ensure coverage, especially for new data.

Curious Replay: Priority Formula

Priority for transition i :

$$p_i \leftarrow c \beta^{v_i} + (|\mathcal{L}_i| + \epsilon)^\alpha$$

- ▶ $c > 0$: base curiosity bonus for rarely trained items
- ▶ $0 < \beta < 1$: decays the curiosity bonus as i is replayed more
- ▶ $\alpha \in (0, 1]$: softens extremes (like in PER)
- ▶ $\epsilon > 0$: avoids zero priority

Sampling probability:

$$P(i) = \frac{p_i}{\sum_j p_j}$$

After sampling i -th transition:

$$v_i \leftarrow v_i + 1$$

Curious Replay: Why It Works

- ▶ **After a change:**
 - ▶ New/shifted transitions have high $|\mathcal{L}_i|$ and low $v_i \rightarrow$ get sampled more.
- ▶ **As learning progresses:**
 - ▶ Both error and curiosity bonus decrease, so focus shifts to other transitions.
- ▶ **Result:** The world model adapts quickly, and the policy can trust its imagined rollouts again.

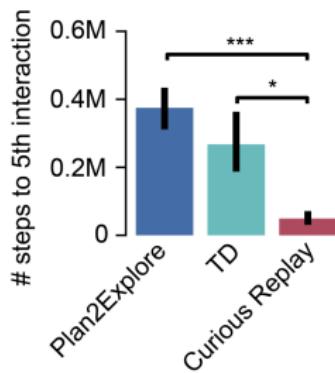
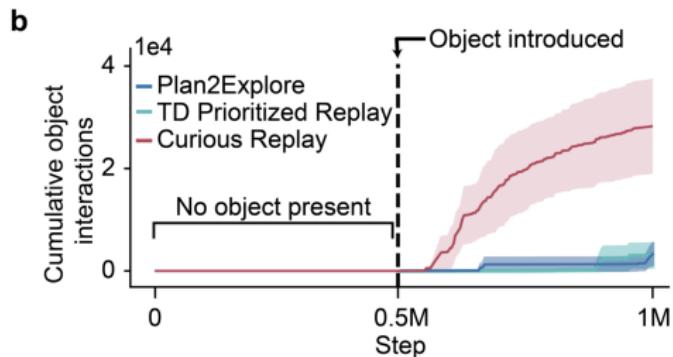
Curious Replay vs. Classic PER

- ▶ **PER:** Prioritizes by value TD error (how surprising the reward is).
- ▶ **Curious Replay:** Prioritizes by world-model error (how surprising the transition is) **and** a freshness bonus.
- ▶ **Why?** In Dreamer, model accuracy is the bottleneck for adaptation after distribution shifts.

Curious Replay: Practical Tips

- ▶ Start with $\alpha \in [0.5, 1]$, $\beta \in [0.8, 0.99]$, small ϵ
- ▶ Choose c so new items get sampled promptly but don't dominate
- ▶ Cap priorities at p_{MAX} to avoid outliers
- ▶ (Optional) Add importance-sampling weights for strict bias control, but the simple scheme works well in practice

Experiment: Interaction with Object Introduced Halfway



Experiment: Constrained Control Released Halfway

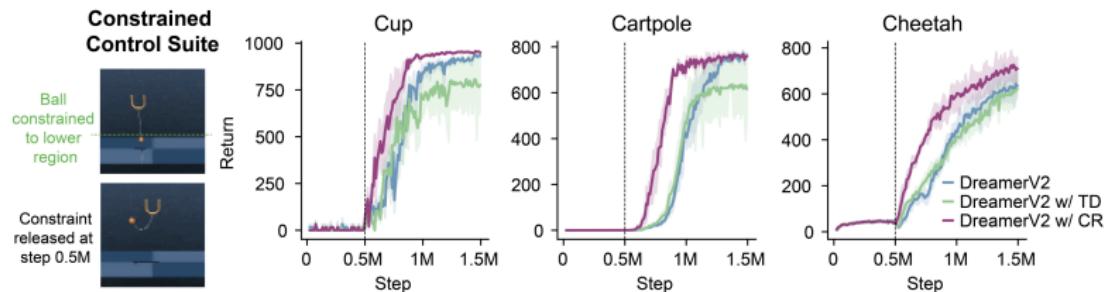


Figure 3. DreamerV2 w/ Curious Replay outperforms DreamerV2 and DreamerV2 w/ TD in the Constrained Control Suite (n=6 per method, mean +/- s.e.m.)

Experiment: CR Helps for Crafter

Method	Crafter Score
DreamerV3 CR	19.4 ± 1.6%
DreamerV3 TD	17.0 ± 2.0%
DreamerV2 CR (8x train freq.)	15.7 ± 2.4%
DreamerV3 [†]	14.5 ± 1.6%
DreamerV2 CR*	13.2 ± 1.4%
LSTM-SPCNN [†]	12.1 ± 0.8%
DreamerV2	11.7 ± 0.5%
DreamerV2 (8x train freq.)	11.0 ± 1.5%
DreamerV2 TD	10.8 ± 0.6%
DreamerV2 [†]	10.0 ± 1.2%
DreamerPro CR	6.8 ± 0.5%
DreamerPro TD	5.8 ± 0.5%
DreamerPro	4.7 ± 0.5%
IRIS	4.6 ± 0.7%
Plan2Explore CR (unsup)	2.7 ± 0.1%
Plan2Explore TD (unsup)	2.7 ± 0.1%
Plan2Explore (unsup)	2.2 ± 0.1%

[†]Published results, see ([Hafner et al., 2023](#))

Experiment: Background Swap in DMC

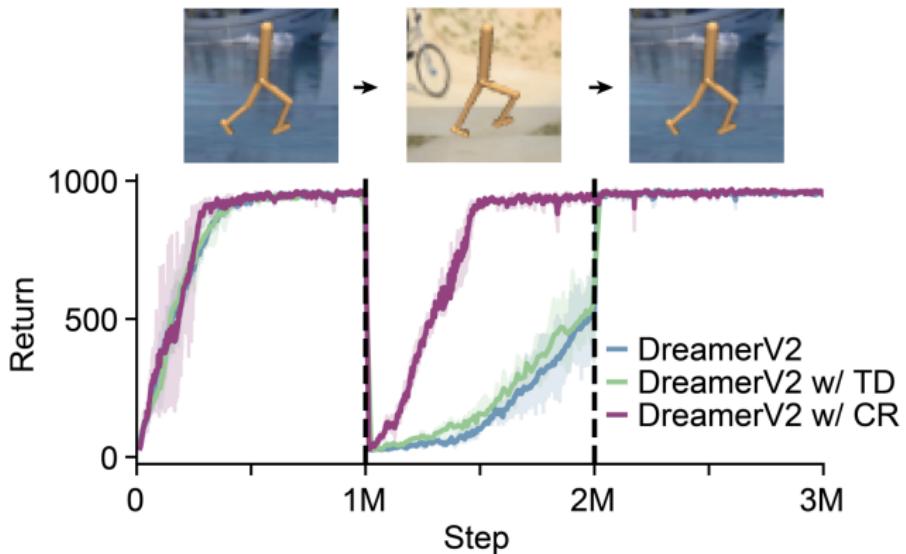


Figure 5. Background-Swap walker_walk. Background changes at step 1M, and reverts at step 2M. Curious Replay improves performance after 1M, without degrading performance after 2M.